

UniMind: Calibration-Aware Scheduling and Placement for Quantum-Accelerated AI Middleware

Y. X. Tan

Abstract—Quantum-accelerated AI workloads must be scheduled onto physical devices whose calibration changes between runs. A classical scheduler assumes resources are identical and stable; a quantum scheduler faces an array of qubits whose individual noise parameters drift daily. This paper asks: *how should a quantum runtime maintain scheduling decisions when hardware calibration changes over time?* We instrument a 156-qubit superconducting device across three daily calibration snapshots (D0–D2; 2026-08-29 to 08-31) and build UniMind, a middleware that scores, places, executes, and refreshes calibration data in a closed loop. Three findings are established: (1) Calibration validity is temporally bounded. Day-over-day rank correlation is only $\rho=0.579$, top-10 overlap is 0.25, and a stale top-3 pin is $2.1\times$ worse than a fresh selection within 24 h. We formalize the **Calibration Validity Horizon** T_{valid} : the time window over which a calibration-derived placement remains within fidelity tolerance. (2) Adaptive refresh beats static and periodic policies. A two-line trigger (top-10 Jaccard < 0.5 or $\Delta C > 0.005$) detects staleness at near-zero cost (~ 0.04 ms per scoring pass). Offline threshold sweeps over $\tau_J \in [0.1, 0.9]$ and $\tau_C \in [0.001, 0.02]$ reveal the Pareto tradeoff between refresh rate and fidelity; the measured trigger sits near the knee. (3) Refresh, not placement selection, dominates end-to-end reliability. Fresh pins lift *both* the pinned and free arms by +33/37 pp on a 6-qubit angle suite; the pinned vs. free difference is within noise (70.0% vs. 76.7%, $p=0.56$). A negative result completes the picture: re-fitting the 7-dim layout utility $J(G)$ on real device labels collapses its out-of-sample Spearman from $\rho=0.818$ to 0.13 ($p=0.36$), because the simulator proxy over-penalizes large layouts that a refresh-pinned device does not. We give the validity model, the measured numbers, the adaptive mechanism, and the explicit negative results.

Index Terms—quantum computing middleware, calibration-aware scheduling, calibration validity horizon, adaptive refresh, qubit placement, calibration drift, resource management, quantum-cloud runtime, reliability composition.

I. INTRODUCTION

EMBEDDED AI is increasingly offloaded to accelerator hardware, and quantum processors are the newest such accelerator. Their software stack, however, changes the *scheduler's* job in a fundamental way: a quantum circuit's fidelity depends on *which* physical qubit executes it and *when* that qubit was last calibrated [1], [2], [5], [6]. A classical scheduler assumes resources are identical and stable; a quantum scheduler faces an array of hundreds of qubits whose individual noise parameters move between runs [10], [13], [17].

Quantum-machine-learning (QML) frameworks—PennyLane [4], Qiskit [18], TensorFlow Quantum—provide algorithm abstractions but not the *system* layer: dynamic backend discovery, layout selection, and calibration refresh.

Y. X. Tan is an independent researcher (e-mail: yxtan5022@gmail.com). All experiments are open-source and reproducible from yxtan5022-maker/unimind-v2; QPU job IDs and snapshots are recorded in the data directory.

Applications are assumed, usually silently, to have been placed on “good enough” qubits. On real devices this assumption fails measurably: with *free* transpiler placement an angle-encoding experiment on `ibm_marrakesh` missed the 0.05 fidelity tolerance by up to 0.392, while the *same* circuits on a pinned calibrated layout passed everywhere (max. deviation 0.028) [25]. Placement is not cosmetic; it is the dominant accuracy term at our test scale.

This paper addresses the temporal dimension of that problem. Placement decisions are made from a calibration snapshot, and snapshots change daily; the scheduler's real questions are “*how old is my snapshot*” and “*should I refresh*” as much as “*where to place*”. We define the **Calibration Validity Horizon** T_{valid} : the maximum age of a calibration snapshot such that the resulting placement remains within fidelity tolerance. With T_{valid} small (hours, not days), the runtime must refresh frequently; with T_{valid} large, refresh is wasteful. We measure T_{valid} 's lower bound on a 156-qubit device across three daily snapshots, design an adaptive refresh trigger that detects staleness at near-zero cost, and show that refresh—not placement selection—is where reliability engineering should invest. The middle section of the paper is a deliberately adversarial audit of our own utility function, because the positive result (§VII, §VIII) is precisely the one the optimistic number (proxy-trained $\rho=0.82$) would have hidden.

A. Contributions

All claims below are measured on the 156-qubit `ibm_marrakesh` device across three daily calibration snapshots (D0–D2; 2026-08-29 to 08-31). We make three contributions:

- 1) **Temporal validity analysis.** We quantify day-over-day calibration drift ($\rho=0.579$, $\tau=0.431$, top-10 Jaccard 0.25), measure the staleness cost of aged pins (stale top-3 is $2.1\times$ worse than fresh within 24 h), and formalize the Calibration Validity Horizon T_{valid} as a design primitive for quantum runtimes (§VII).
- 2) **Adaptive refresh policy.** We design a trigger (Jaccard < 0.5 or $\Delta C > 0.005$) that detects staleness at near-zero cost, run offline threshold sweeps over τ_J and τ_C to reveal the refresh-rate-vs.fidelity Pareto tradeoff, and compare static, periodic, and adaptive policies on the same data (§VII, §VIII).
- 3) **Negative results.** The 7-dim layout utility $J(G)$, trained on a simulator proxy and re-fitted on real device labels, collapses from out-of-sample $\rho=0.818$ to 0.13 ($p=0.36$); at $k=6$ with fresh pins, pinned placement shows no measured advantage over the free transpiler (70.0% vs.

TABLE I
NOTATION

Symbol	Meaning
S_t	calibration snapshot at time t (readout errors, gate errors, T_1, T_2)
$C(q)$	total readout assignment error $p_{0 1}(q) + p_{1 0}(q)$
ρ, τ	Spearman, Kendall rank correlation across $S_t \rightarrow S_{t'}$
k	number of logical qubits of a circuit
$\pi : \hat{G} \rightarrow Q$	placement of the logical circuit on physical qubits
w	amplitude weight of one encoded value; $\theta = 2 \arcsin \sqrt{w}$
$Z_{\text{emp}}, Z_{\text{th}}$	empirical, theoretical Z -measurement
ϵ	fidelity tolerance (0.05); pass iff $\max_i Z_i^{\text{emp}} - Z_i^{\text{th}} \leq \epsilon$
$P_{\text{obs}}(1) = b + aw$	two-parameter affine readout channel model
$J(G)$	7-dim layout utility with weights $(\alpha, \beta, \gamma, \gamma_{\text{log}}, \eta_{\text{idle}}, \delta, \lambda)$
N_{SWAP}, D	inserted SWAPs and transpiled depth
g, r, f	per-call orchestrator outcome probs. (correct, broken, unavailable)
H, K	healing budget, generation-call budget
$R_{\text{end}}, \rho_{\text{heal}}, C$	end-to-end reliability, healed-path reliability, fallback coverage

76.7%, $p=0.56$). We report these as bounds on what placement engineering can deliver without temporal control (§VI, §VIII).

Scope. We make no claim of quantum advantage, novel quantum algorithms, or a production runtime. We propose T_{valid} as a design primitive; with three daily snapshots we establish its lower bound and demonstrate the adaptive mechanism, but leave the full multi-day decay curve and cross-device generalization to future work. Every quantitative claim traces to a recorded job or snapshot (§X); where an exploratory result did not replicate (a first run’s perfect rank correlation, $\rho=1.000$, which a fixed-design replication returned as $\rho=0.771$, $p=0.103$) we report it with its failure. UniMind also contains an LLM-orchestration layer whose reliability composition we *do* model exactly (§VIII-A), but the LLM’s quality is a controlled parameter of our experiments, never a claim about real models. **Organization.** Section III formalizes the scheduling problem and surveys related work; §IV describes the middleware; §V–VI establish the placement results and their limits; §VII measures drift and validates the refresh policy; §VIII composes the software and hardware paths; §IX states what the measurements do and do not say. Table I fixes notation.

II. EXPERIMENTAL SETUP

A. Device and snapshots

All hardware results use the IBM Quantum 156-qubit superconducting device `ibm_marrakesh` (heavy-hex topology, 176 couplers). We captured three full calibration snapshots—D0, D1, D2—at roughly 24-hour and 10-hour spacing and ran every device experiment against the matching snapshot

TABLE II
SNAPSHOT TIMELINE. PINS FOR EACH EXPERIMENT FAMILY REFERENCE THE STATED SNAPSHOT; EXECUTION TIMES ARE THE JOB QUEUE WINDOWS RECORDED IN THE DATA FILES.

Label	device update	capture	used as pins for
legacy 08-23	2026-08-23	2026-08-23	§V anchor
D0	2026-08-29 21:26	22:07	§V-B, stale arms
D1	2026-08-30 21:29	23:03	drift/staleness eval
D2	2026-08-31 07:46	—	refresh arms (§VI, §VIII)

(Table II). Two snapshot families recur: the *pinned* family `calib_2026-08-{29, 30, 31}.json`, and a legacy 08-23 capture used only for the placement anchor of §V. Where “stale” pins are evaluated, the pins are always taken from D0 while execution happens on D1/D2-era device state.

B. Common protocol

The instrument circuit is angle (RY) encoding on Z -measurement, 19-weight grid $w=0.05, \dots, 0.95$ for the single-qubit sweeps and 30 parameterizations for the 6-qubit batches; transpiler seed 42 throughout; layouts are pinned via `initial_layout` where placement is under test and left free where it is the baseline. Shots are 8192 (single-qubit sweeps) or 4096 (multi-qubit suites). The pass criterion is per-bit $\max_i |Z_i^{\text{emp}} - Z_i^{\text{th}}| \leq \epsilon$ with $\epsilon=0.05$, fixed before data release. All statistical tests are two-sided unless stated; proportions use Wilson intervals and two-proportion z -tests computed from the raw pass counts. Details that would be tedious in the main text (per-row values) are in the appendix.

III. PROBLEM AND RELATED WORK

A. Formal model

Definition 1 (Calibration-aware scheduling problem): A device is a set of qubits Q and couplers E with a calibration snapshot S_t reporting, per qubit, readout assignment errors ($p_{0|1}, p_{1|0}$), gate errors, and coherence times (T_1, T_2). A job is a logical circuit $\hat{G}$ over k logical qubits. The scheduler chooses a placement $\pi : \hat{G} \rightarrow Q$ (with $|\pi| = k$, connected in (Q, E) for our chains), a transpiler configuration, and a *snapshot age policy* that decides when S_t is re-fetched. Its objective, subject to placement, snapshot age, and overhead budgets, is per-qubit measurement fidelity $\max_i |Z_i^{\text{emp}} - Z_i^{\text{th}}| \leq \epsilon$ at a reliability target.

The distinctness from classical placement lies in (Q, E, S_t) being jointly non-stationary: π chosen at time t is evaluated at execution time $t' > t$ against a changed $S_{t'}$. The paper’s arithmetic is about the size of that violation. Three assumptions bound the model.

- **(A1) Readout dominates measurement.** For the Z -measurement instrument used throughout, per-qubit readout assignment error is the leading distortion term; gate noise is handled by the placement/transpiler and by coupler terms in $J(G)$, but readout is the mechanism we isolate first.
- **(A2) Snapshots are the operative truth.** The scheduler sees the device only through calibration snapshots; it cannot poll the device state at shot granularity.

- **(A3) Refresh is cheap.** A snapshot fetch and a full-156 scoring pass cost ≈ 0.04 – 0.4 ms and one API call; under A2, staleness is an *operational-policy* failure, not a compute one.

Definition 2 (Calibration Validity Horizon): Let π_t denote a placement computed from snapshot S_t , and let $F(\pi_t, S_{t'})$ denote the execution fidelity of π_t evaluated on device state $S_{t'}$ at time $t' > t$. For a fidelity tolerance $\epsilon > 0$, the *Calibration Validity Horizon* is

$$T_{\text{valid}}(\pi_t, \epsilon) = \max\{\tau \geq 0 : F(\pi_t, S_{t+\tau}) \geq F(\pi_t, S_t) - \epsilon\}. \quad (1)$$

T_{valid} depends on device, workload, k , and metric; it is not a universal constant.

B. Related work

Placement and compilation. Qubit mapping has both heuristic and exact treatments: SABRE [6] and noise-adaptive mappings [5] fold device error into the routing pass; exact synthesis is practical for small circuits [7]; resource-aware timing variants exist for heterogeneous devices [8]. Our contribution is not a new routing algorithm but *who decides*: a middleware layer that computes an `initial_layout` from a score over a live snapshot and composes with any downstream transpiler, plus an explicit staleness contract—aspects the compiler literature treats as “the backend’s problem”.

Error mitigation. Measurement twirling [11], [12], probabilistic error cancellation, and sparse Pauli–Lindblad models [9], [10] correct errors statistically at shot cost. Our results bound where mitigation is needed: it helps the *broken* qubit and is redundant on the good one (§V), which argues for calibration-conditional mitigation—a placement decision, not a post-processing one. Dynamical decoupling [15] is orthogonal and composes with our layout.

Runtime and cloud middleware. Qiskit Runtime [16], Braket [20], and Azure Quantum [21] manage jobs and queue bounds but expose placement as an option, not as a scheduler resource; much of the practical “does this job need mitigation, where, and with which snapshot” knowledge lives in operator runbooks. The transpiler layer itself (III-B) exposes `initial_layout` but no policy over it. Table III summarizes the capability split. We know of no published quantified day-scale ranking drift for a 156-qubit device; the public signals closest to it are per-device aggregate fidelities in documentation [17] and in error-rate studies of specific devices [3]. We present the first measurements of the drift and the protocol (trigger, refit-on-real-labels) that a runtime can adopt.

Calibration data as a first-class resource. In the classical substrate the analog is fast-varying per-node health signals (e.g. network or DRAM ECC counters) that schedulers already timestamp and age [8]. Device vendors publish aggregate fidelities per backend and expose per-qubit snapshots through an API ([17]), i.e. the *data* for a calibration-aware scheduler exists at negligible cost; the missing piece is a runtime policy that (a) validates the score’s predictive content and (b) bounds its validity horizon explicitly—both of which this paper measures. We stress that the relevant comparison class is

TABLE III
CAPABILITY SPLIT BETWEEN QUANTUM SOFTWARE LAYERS AND UNIMIND.
 $\checkmark/\times/\cdot$ = SUPPORTED / ABSENT / IMPLICIT-ONLY.

Layer	ALG.	LAYOUT	SNAP.	REFRESH	RELIAB.
Qiskit/PennyLane [4], [18]	$\checkmark$	$\times$	$\cdot$	$\times$	$\times$
Qiskit transpiler [6]	$\times$	$\checkmark$	$\cdot$	$\times$	$\times$
Qiskit Runtime [16]	$\checkmark$	$\cdot$	$\cdot$	$\times$	$\cdot$
Braket [20]/Azure [21]	$\checkmark$	$\cdot$	$\cdot$	$\times$	$\cdot$
Mitiq [14]/Qiskit-mit [11]	$\cdot$	$\times$	$\cdot$	$\times$	$\checkmark$
CUDA-Q [19]	$\checkmark$	$\cdot$	$\cdot$	$\times$	$\times$
UniMind (this paper)	$\checkmark$	$\checkmark$	$\checkmark$	$\checkmark$	$\checkmark$




Fig. 1. The UniMind scheduling loop. One scoring pass is sub-millisecond; the refresh trigger decides when the stored snapshot must be re-fetched. (execution $\rightarrow$ results) is emitted downstream.

not “our scheduler beats vendor plumbing” but “a score whose age is policed beats the same score whose age is not,” which is the controlled experiment of §VIII.

IV. UNIMIND MIDDLEWARE

UniMind is a user-space, multi-backend middleware (Qiskit/Aer, CUDA-Q [19] via Qiskit, IBM Quantum). Its scheduling core is the calibration-aware loop of Fig. 1: one snapshot per calibration cycle, $C(q)$ scoring, chain/utility placement, execution, and a refresh trigger. Three layers are described in turn: the *Unibit* preprocessing layer that turns a classical amplitude sequence into encoding parameters; the *orchestrator* that interprets task intent and enforces a hierarchical fallback; and the *placement core* (score, chain, utility, trigger) that the previous two sections measure.

A. Unibit preprocessing

From a classical amplitude sequence $x_0, \dots, x_{n-1}$, Unibit derives the w_i encoded by each gate: a sliding-window frequency estimator (window width 3–5 samples) with sinc-shaped smoothing and a fixed phase shift $\phi=0.42$ that breaks sampling-grid symmetry. The encoding is the standard angle (RY) form [22], [23], $\theta_i = 2 \arcsin \sqrt{w_i}$, which we use strictly as an *instrument* to expose placement and drift—the same circuits everywhere in the paper. The preprocessing step carries 13 formal invariants (kernel positivity and monotonicity, adaptive-threshold equivalence, dead-zone boundary matching the sinc root, shot-level consistency, affine-channel closure, etc.), all 13 verified by property tests against the implementation; none is claimed as a novelty, and the derivation is open-source. What matters for the scheduler is that $w \mapsto \theta$ is deterministic and

that the Z -basis measurement is the mechanism the rest of the paper disturbs by calibration.

B. Oracle/orchestrator layer

The *orchestrator* interprets a task intent into candidate code under sandboxed validation (static analysis, pattern checks, semantic checks) with a deterministic rule-based fallback, so every intent has a plan even when generation fails—this hierarchy is what §VIII-A models exactly. Its two-parameter quality model (q = per-call correctness, f = unavailability) is a controlled experimental parameter in this paper; the LLM itself is never part of a hardware claim. The software side is composed with the hardware side in §VIII.

C. Placement core

Every qubit is scored by total readout assignment error

$$C(q) = p_{0|1}(q) + p_{1|0}(q), \quad (2)$$

ties broken by $-\min(T_1, T_2)$. The score is fixed *a priori*, never fitted; it is tested out-of-sample across the entire calibration spectrum (§V-B) rather than tuned on the data it later predicts.

For a k -qubit circuit we grow a *connected chain*: from the best $C(q)$ qubit, add the lowest- C neighbour of the current set (greedy Prim-style expansion; 0.3 ms at $k=3$, scales to the full device); an exact BFS frontier search is used for the counted suites of §VI. For utility-based ranking among candidate layouts,

$$J(G) = \alpha E_{\text{readout}} + \beta E_{1q} + \gamma E_{2q} + \gamma_{\log} E_{2q_{\log}} + \eta_{\text{idle}} E_{\text{idle}} + \delta N_{\text{SWAP}} + \lambda D, \quad (3)$$

with E_{readout} the mean $C(q)$ of placed qubits, E_{1q} mean single-qubit gate error, $E_{2q} = \sum_c p_{cz}(c)$, $E_{2q_{\log}} = \sum_c -\log(1 - p_{cz}(c))$ (log-odds coupler cost), $E_{\text{idle}} = \text{depth} \cdot \text{mean}(1/T_1 + 1/T_2)$, N_{SWAP} inserted SWAPs, and D transpiled depth; features min–max normalized over training points. Section VI measures what this utility is worth when its labels come from the real device. For interactive routing we also support calibration-weighted SABRE: node weights $w(q) = 1/(C(q) + 3 s_x(q))$ seed the SABRE candidate set, the best of the weighted candidates and the global SABRE layout is kept (so $k \geq 8$ is never worse than default), and a guarantee check verifies zero violations. In the measured comparison at $k = 2, 4, 6, 8, 10, 16$ this removes 76–100% of the SWAPs of a random baseline while matching default depth exactly.

D. Refresh trigger

After each calibration cycle the middleware recomputes the top-10 of $C(q)$ and refreshes the stored snapshot when (i) top-10 Jaccard overlap with the stored snapshot falls below 0.5, or (ii) the stored best qubit’s C rises by more than 0.005 absolute. Re-scoring is ~ 0.04 ms—economically free—so staleness is purely an operational-policy failure, quantified next.

TABLE IV
PLACEMENT-CONTROLLED SWEEP (2026-08-23; 19 WEIGHTS, 8192 SHOTS, 3 JOBS/CELL). DEVIATIONS ARE $\max_w |Z_{\text{emp}} - Z_{\text{th}}|$; TOLERANCE 0.05.

Placement	$C(q)$	Bare	Twirl	Verdict
q98 (best)	0.0044	0.028	0.031	pass
q0 (free)	—	0.026	—	pass
q37 (adversarial)	0.82	0.213	0.144	fail
v1.7 free layout	unrec.	0.392	0.245	fail

E. System guarantees

UniMind enforces three invariants aimed at minimizing regressions against a baseline transpiler. First, *non-degradation*: in calibration-weighted SABRE, the chosen layout is always the better of the weighted-SABRE candidate and the global-SABRE default; the guarantee check ($k=2, 4, 6, 8, 10, 16$) reports zero violations against depth and SWAP counts. Second, *deterministic fallback*: when the oracle generates `None`, a rule-based fallback compiles a template; when the template itself fails (e.g. apostrophes breaking interpolation), a last-resort closure produces the best available partial output, preserving completeness at the cost of fidelity (this is exactly what the $c=2/9$ fallback coverage of §VIII-A measures). Third, *reproducibility*: transpiler seed is fixed globally (42), layout pins are snapshot-addressed, and every job ID is returned to the caller so the execution trace can be recovered from the vendor API. These guarantees do not require modifying the quantum programming model and can be overlaid on any backend that exposes an `initial_layout` parameter and a calibration-data endpoint.

V. PLACEMENT-DOMINATED FIDELITY ON ONE QUBIT

The anchor experiment is a 2×2 design: {best-calibrated qubit, adversarially bad qubit} $\times$ {bare, measure-twirled}, on the 19-weight grid $w \in \{0.05, \dots, 0.95\}$ of the amplitude encoding, 8192 shots, transpiler seed 42, layouts pinned via `initial_layout`, 3 jobs per cell. On the 2026-08-23 snapshot, q98 minimizes C at 0.0044 ($T_1=207 \mu\text{s}$, $T_2=67 \mu\text{s}$; q37 is the adversarial anchor ($C=0.82$)).

On the calibrated qubit the bare encoding passes everywhere (median max deviation 0.028; job range 0.020–0.029) and the transfer curve is fit at shot-noise level by $P_{\text{obs}}(1) = 0.004 + 0.989 w$ (residual RMS 0.0042, below the ≈ 0.0055 shot-noise floor), with per-repeat slopes 0.986–0.994. On q37 the same circuits fail by $\sim 4\times$, and the fitted channel acquires a large offset, $P_{\text{obs}} = 0.121 + 0.855 w$ —the signature of an asymmetric readout channel. Twirling changes nothing significant on the good qubit (slopes 0.978–0.987, RMS 0.005–0.008) and only partially rescues the bad one ($0.213 \rightarrow 0.144$). Fig. 2 plots the transfer curves; Table V collects the per-cell fits. Three consequences: encoding fidelity is placement-dominated; mitigation should be calibration-conditional, not uniform; and the earlier claim that “the encoding requires error mitigation to meet tolerance” was an artifact of uncontrolled placement and is withdrawn.

TABLE V

AFFINE CHANNEL FITS $P_{\text{obs}} = b + a w$ ACROSS THE PLACEMENT DESIGN (MEDIAN OF 3 JOBS/CELL). RMS IS ON P_{obs} ; SHOT-NOISE FLOOR ≈ 0.0055 .

Cell	slope a	offset b	RMS	max dev
q98 bare pinned	0.989	0.004	0.0044	0.028
q98 twirled	0.987	0.011	0.0066	0.031
q0 free	0.993	-0.002	0.0067	0.026
Aer local model	1.007	-0.003	0.0030	0.015
q37 bare pinned	0.855	0.121	—	0.213




Fig. 2. Placement-dominated encoding fidelity: the bare response on the best-calibrated qubit tracks the ideal line within shot noise; the adversarially chosen q37 compresses the curve toward an offset set by its asymmetric readout.

A. Calibration predictability

Across the whole design, at shot-noise level,

$$P_{\text{obs}}(1) = b + a w, \quad Z_{\text{obs}} = d + c Z_{\text{th}}, \quad c=a, \quad d=1-2b-a, \quad (4)$$

with $(a, b) = (0.99, 0.00)$ on the good qubit vs. $(0.86, 0.12)$ on the bad one; the offset matches the direction and scale of that qubit’s readout asymmetry. The one-parameter contraction model $P_{\text{obs}} = 0.5 + \alpha(w - 0.5)$ is the balanced-readout special case $b = (1 - a)/2$ and is rejected where asymmetry dominates. Matched calibration–noise-model simulation on Aer explains most of observed error on the good qubit (median QPU/model deviation ratio $1.4\times$, Table V): the “standard models underestimate hardware” gap is largely a placement confound, not a model gap. The affine model is also closed under composition (the composed channel of two affine distortions is affine with slope product and offset rule verified by property tests), which is why Eq. (4) remains usable after cascaded stages.

B. Stratified validation

From the D0 snapshot all 156 qubits are ranked: q98 is 1/156 ($C=0.004$), q37 is 128/156 ($C=0.10$). Six qubits sampled across the ranking run the identical bare 19-weight sweep (Table VI). The top qubit passes all 19 weights (0.030 max dev; affine slope degrades along the ranking in the predicted direction: a : 0.986 $\rightarrow$ 0.887, offset b : +0.001 $\rightarrow$ +0.081), but the trend is *not monotone*: q109 at percentile 25.2% already fails, so Spearman $\rho=0.771$ at $n=6$ is not significant (exact permutation $p=0.103$) and a “top-half rule” is unsupported. An earlier exploratory run (first snapshot, 08-22) was perfectly monotone ($\rho=1.000$, minimal attainable p); that result did not survive the fixed-design replication and is reported as

TABLE VI

$C(q)$ -STRATIFIED SWEEP (D0 SNAPSHOT; SIX QUBITS $\times$ 19 WEIGHTS $\times$ 8192 SHOTS). ‘PASS’ = MAX DEVIATION < 0.05 OVER ALL 19 WEIGHTS; (a, b) FROM EQ. (4).

Qubit	Rank pct.	$C(q)$	max dev	pass	(a, b)
q98	0.0%	0.004	0.030	19/19	(0.986, +0.001)
q109	25.2%	0.017	0.081	14/19	(0.933, +0.043)
q105	50.3%	0.031	0.043	19/19	(0.983, +0.018)
q53	75.5%	0.072	0.046	19/19	(0.982, +0.021)
q37	81.9%	0.099	0.146	9/19	(0.887, +0.081)
q27	95.5%	0.346	0.096	13/19	(0.917, +0.049)




Fig. 3. $C(q)$ -guided ranking validation: top-ranked qubits pass, but deviations are not monotone along the ranking (q109 at percentile 25.2% fails), and score vs. deviation gives Spearman 0.771 ($p=0.103$, $n=6$).

unreplicated. Overhead: 0.036 ms full-156 ranking; pinned transpilation 1.9 ms vs. 1.8 ms free.

C. Multi-qubit routing

Single-qubit $C(q)$ is necessary but not sufficient. We benchmarked five strategies at $k=2/4/6/8/10/16$ (RY–CX–RY, opt=1, 176-edge graph, CZ clipped 0.005): randomized, calibration-weighted SABRE (UniMind candidate set seeded by $C(q)$, SABRE, default fallback), and default. UniMind reduces SWAPs by 76–100% vs. random and matches default for $k \geq 8$ by fallback (the fallback keeps the *better* of the weighted-SABRE and global-SABRE layouts; guarantee check reports zero violations); the log-odds coupler cost separates high-error couplers (Random 1.72 vs. UniMind 0.14 at $k=8$). Latency stays below 20 ms (pick < 1 ms, transpile dominated). This is a *system* result about routing cost; whether the chosen layout is *correct* is what §VI measures.

VI. THE $J(G)$ UTILITY: PROXY FIT VS. REAL LABELS

A. Fit protocol

The seven weights of Eq. (3) are fitted by max-Spearman simplex search (Dirichlet+corners sampler, 5 000 samples, seed 42) on $n=23$ points: six single-qubit real max-dev labels from §V-B and seventeen multi-qubit points $k=2-18$ whose label is the simulated two-qubit error E_{2q} of the placed layout under the device error model. Validation: leave-one-out (LOO) and a bootstrap ($B=200$). A *size out-of-sample* split additionally holds out the entire $k \geq 9$ range: fit on $k \leq 8$ only, evaluate on $k \geq 9$ with train-only min–max scaling. All numbers in Table VII use the identical feature matrix; only the labels differ—the vulnerable object is the *label*, not the features. The reason we

TABLE VII

FITTING $J(G)$ ON SIMULATOR-PROXY VS. REAL-DEVICE LABELS. SAME FEATURE MATRIX AND PROTOCOL; THE $k \geq 9$ LABELS DIFFER. SIZE-OOS = FIT ON $k \leq 8$, EVALUATE $k \geq 9$.

Labels	full fit	LOO	size-OOS
v4 proxy ($k=9-18$ est-2q)	0.967	0.953	0.818 ($p=0.0029$)
v5 real (7 measured k)	0.643	0.510	0.127 ($p=0.36$)
Δ	-0.324	-0.443	-0.691

TABLE VIII

MEASURED PER- k MAX-DEV (MEAN OF TWO PATTERNS) OF THE SAME 14-CIRCUIT $k=9-18$ SUITE UNDER STALE (D0) VS. FRESH (D1/D2) PINS, 4096 SHOTS. SPEARMAN k VS. MAX-DEV: STALE $\rho=0.62$ ($p=0.018$), FRESH $\rho=0.15$ ($p=0.61$).

k	9	10	11	12	14	16	18
stale pins	.098	.103	.097	.087	.139	.252	.241
fresh pins	.060	.044	.062	.040	.057	.063	.068

can separate the two is that for $k \in \{9, 10, 11, 12, 14, 16, 18\}$ we reran the same greedy-chain circuits on the device with refresh-pinned layouts and replaced the proxy labels with the measured mean max deviation (Table VIII), giving the v5 dataset.

B. Ranking objective

$J(G)$ is scored by *ordering*, not calibration: the fitted parameters maximize the Spearman correlation between J and the label vector, which is the quantity a dispatcher actually consumes (it picks the argmax over candidate layouts, rarely the calibrated value). This makes the objective conservative in the paper's favor: a model that merely preserved ranks would score $\rho=1$ even with badly calibrated magnitudes, so any measured $\rho < 0.8$ is a genuine ordering failure, not a scale artifact. All fits use min-max-normalized features on the *training* partition only ($k \leq 8$ in the split), so out-of-sample scores $\rho=0.13$ cannot be charged to label leakage; the LOO spread of train rho (0.60–0.72 across folds, Table VII data) also brackets the full-fit 0.643 and shows the v5 weights are stable under 22/23 training sets—the collapse concentrates in the size direction, as the split test isolates.

C. Result

Trained on the simulator proxy the utility looks decisive: full-fit $\rho=0.967$, LOO 0.953, bootstrap 0.97 ± 0.02 , and the size-OOS split $\rho=0.818$ ($p=0.0029$)—holding out entire sizes and still ranking them. The proxy's assumption failed where it mattered. A 14-circuit suite ($k \in \{9, 10, 11, 12, 14, 16, 18\}$, two seeded patterns each, greedy-chain placement) measured real max-dev with *fresh* pins: per- k means are 0.040–0.068 with *no* size scaling (7/14 pass 0.05), while the identical suite pinned from D0 (two days stale) climbs to 0.24–0.25 at $k=16, 18$ with only 2/14 passing (Table VIII, Fig. 4). The stale population's failures concentrate on late-chain qubits (q37, q45–q47, q57) that the greedy policy must include for $k \geq 11$ (the appendix table gives the 14 rows); per-bit $C(q)$ vs.




Fig. 4. The size scaling of the $k=9-18$ suite is an artifact of stale pins: the same circuits pinned from a fresh snapshot are approximately k -flat, which is exactly the premise the simulator proxy got wrong.

measured deviation correlates only $\rho=0.21$. The mechanism is now explicit: simulated E_{2q} keeps rising with k (its toxicity is dose-dependent on k), whereas a refresh-pinned device is approximately k -flat, so the proxy-trained utility over-penalizes large layouts that are in fact fine. Re-fitting on the seven real labels replaces in-sample $\rho=0.967$ with $\rho=0.643$ and destroys the size-OOS claim ($\rho=0.127$, $p=0.36$)—the apparent generalization was a label artifact (Table VII). The learned weights move accordingly (Table X): the proxy fit leans on coupler costs ($\gamma=0.37$), the real-label fit on single-qubit terms ($\beta=0.42$) and SWAP/depth ($\delta=0.20$, $\lambda=0.03$), i.e. the model's opinion of *what matters* flips. The constructive reading is what we adopt in the design: $J(G)$ is usable as a *ranking* device only when trained on real, refresh-pinned labels and applied within a calibration cycle; weights from proxy datasets must not be extrapolated to $k \geq 11$; and chain placement's forced inclusion of boundary qubits is the term to attack next.

Which term survives contact with the hardware? A leave-one-term-out ablation on the real-label OOS split answers this. Refitting $J(G)$ with one interaction dropped and measuring the size-OOS ρ pins the blame (Table IX): removing the coupler term E_{2q} (whose γ is the layout penalty against boundary qubits) collapses OOS ρ from 0.127 to -0.042 ($\Delta\rho = -0.170$), the only drop that changes the sign of the result. Removing E_{1q} costs 0.024; removing the readout, idle, or interaction-density terms costs 0.000; and removing N_{SWAP} actually *raises* OOS to 0.139, indicting it as a pure proxy artifact with no real-label support. The single real signal in $J(G)$ —and the term a hardware-informed design must keep—is the coupler/layout term that encodes where on the device the chain is pinned, which is precisely the information a refresh-driven scheduler can provide.

VII. DRIFT DYNAMICS AND THE REFRESH POLICY

This section measures how fast the score goes stale on two full 156-qubit snapshots D0 (update 2026-08-29 21:26) and D1 (2026-08-30 21:29), with a third capture D2 (2026-08-31 07:46).

A. Global drift

Day-over-day ranking stability is moderate: Spearman $\rho=0.579$, Kendall $\tau=0.431$, top-10 Jaccard 0.25 (only four of D0's best ten survive). The best qubit turns over: q8 ($C=0.0039$, rank 1) rises to 0.0095 (rank 6) while q19

TABLE IX

LEAVE-ONE-TERM-OUT ABLATION OF $J(G)$ ON THE REAL-LABEL SIZE-OOS SPLIT. $\Delta\rho$ IS THE CHANGE IN SPEARMAN ρ RELATIVE TO THE FULL MODEL (0.127). REMOVING E_{2q} IS THE ONLY DROP THAT REVERSES THE SIGN.

Dropped term	OOS ρ	$\Delta\rho$
none (full)	0.127	—
coupler E_{2q}	-0.042	-0.170
single-qubit E_{1q}	0.103	-0.024
readout E_{readout}	0.127	0.000
idle E_{idle}	0.127	0.000
interaction D	0.127	0.000
SWAP N_{SWAP}	0.139	+0.012

TABLE X

LEARNED WEIGHTS OF $J(G)$ UNDER PROXY AND REAL LABELS. OOS-TRAIN = WEIGHTS FITTED ON $k \leq 8$ ONLY (V5 SIZE-OOS SPLIT).

Fit	α	β	γ	γ_ℓ	η_{idle}	δ	λ
v4 proxy (full)	.034	.335	.371	.085	.070	.095	.011
v5 real (full)	.229	.415	.021	.079	.027	.196	.033
v5 real OOS-train	.006	.559	.051	.104	.007	.269	.004

(0.0066) becomes best. Per-qubit spread is extreme (Fig. 5): q37 swings $0.099 \rightarrow 0.740$ (max $|\Delta C| = 0.641$), q53 improves sixfold ($0.072 \rightarrow 0.011$); median $|\Delta C|$ is 0.0127, p90 is 0.123, and 72.4% of qubits move by more than 0.005 (53.9% by 0.01, 41.0% by 0.02). On D2 the top of the ranking is unchanged (best remains q19 with the same top-3 values), consistent with a regime dominated by daily calibration cycles. Table XI tracks the six qubits used in §V-B; note the anticorrelated moves (q53 and q27 improve while q37 worsens), and that rank movement is far larger than C movement would predict near the top—the ordering is fragile exactly where selection operates.

B. Staleness cost and the trigger

Pinning three qubits from the *one-day-old* ranking and evaluating at D1 gives $C = 0.0095/0.0115/0.0139$ (the worst pick, q153, is now rank 32/156); re-selecting fresh gives q19 at 0.0066—the stale policy’s worst pin is $2.1\times$ worse than the adaptable choice (Table XIII). The trigger of §IV-D fires on $D0 \rightarrow D1$ (Jaccard $0.25 < 0.5$; stored-best $\Delta C = 0.0056 > 0.005$) and would not fire $D1 \rightarrow D2$ (top-3 unchanged), matching the regime hypothesis. Because re-selection is sub-millisecond, staleness is a pure operations failure, not a compute constraint.

These two measurements bound T_{valid} (Definition 2) from both sides. For the $k=6$ chain and tolerance $\epsilon=0.05$, the selected pins remain within tolerance on *this device* at $t+10$ h ($D1 \rightarrow D2$ top-3 unchanged, $T_{\text{valid}} \gtrsim 10$ h) but the cost of a one-day-old ranking already reaches $2.1\times$ ($D0 \rightarrow D1$, $T_{\text{valid}} \lesssim 24$ h). T_{valid} is therefore *hours*, not *days*, on this device at this k —small enough that refresh must be part of the runtime loop, not a periodic maintenance step. The exact position between the two bounds is a function of workload, k , and device family (1).




Fig. 5. Per-qubit readout error $C(q)$ at D1 versus D0 (log-log). Qubits swing in both directions: the best qubit at D0, q8, falls out of the top six while q37 worsens by a factor of 7.5. The ranking is informative but not persistent.

TABLE XI

DAY-OVER-DAY MOVEMENT OF THE SIX ROUTED QUBITS OF TABLE VI. RANKS ARE 1–156.

q	$C@D0$	rank@D0	$C@D1$	rank@D1	ΔC	ratio
q8	.0039	1	.0095	6	+0.0056	$2.44\times$
q109	.0171	41	.0154	44	-.0017	$0.90\times$
q105	.0313	80	.0222	68	-.0090	$0.71\times$
q53	.0718	118	.0115	18	-.0603	$0.16\times$
q37	.0991	128	.7397	152	+0.6406	$7.46\times$
q27	.3462	149	.0620	119	-.2842	$0.18\times$

C. Ranking margin and turnover

A sharper mechanism links drift to turnover. Let $M_k = C_{(k+1)} - C_{(k)}$ be the score gap between adjacent ranks: how much the device must drift before the ordering can flip. Because the top of the ranking is packed ($M_1 = 0.0020$, $M_2 = 0.0015$ at D0—separations far smaller than the median daily drift $|\Delta C| = 0.0127$), a single day’s drift exceeds the margin that separates the top candidates, and the best qubit falls (Table XII). Across the top-50 selection-relevant qubits the Spearman correlation between $M_k(D0)$ and the qubit’s rank displacement at D1 is $\rho = -0.425$ —smaller margins tend to predict larger turnover, though the sample ($n=50$, one calibration epoch) is too small to establish significance ($p=0.26$). The effect is local to the top: over all 156 qubits it vanishes ($\rho = -0.09$, n.s.), because far from the top both margins are large and drift proportionally small. The pairing of the two observations explains the regime difference: $D0 \rightarrow D1$ has large drift relative to small top margins, so turnover; at D1 the top margins are even smaller yet the $D1 \rightarrow D2$ drift is near zero, so the ranking is stable despite the packing. Margin is thus the quantity that converts a drift magnitude into a selection failure: the runtime should refresh not on fixed wall-clock but when the *ratio* of accumulated drift to the current top margin approaches unity.

D. Direct device test

The refresh-pinned suite of §VI is the direct demonstration: the *same* circuits, differing only in pin freshness, measure means 0.04–0.07 (fresh pins, 7/14 pass) vs. 0.10–0.25 (stale, 2/14 pass), and the large- k degradation disappears entirely (Table VIII, Fig. 4). The dominant failing qubit under staleness is q37 ($C=0.74$ at D1). Refresh—not more elaborate selection—is what keeps errors bounded at $k \geq 9$.

TABLE XII

WHAT HAPPENS TO SELECTED D0-RANKED QUBITS WHEN A SINGLE DAY OF CALIBRATION DRIFT ARRIVES. “MARGIN” IS THE D0 SCORE GAP BELOW THE QUBIT; THE DRIFT $\rightarrow$ D1 THAT EXCEEDS IT REORDERS THE RANKING. (RANKS 6, 9–11 ARE OMITTED FOR BREVITY.)

rank@D0	q	Margin	C@D0	C@D1	ΔC	rank@D1
1	q8	.0020	.0039	.0095	+.0056	6
2	q120	.0015	.0059	.0115	+.0056	19
3	q153	.0005	.0073	.0139	+.0066	32
4	q4	.0005	.0078	.0173	+.0095	52
5	q7	.0000	.0083	.0244	+.0161	78
7	q154	.0005	.0083	.0071	−.0012	3
8	q6	.0005	.0088	.0073	−.0015	4
12	q19	.0000	.0098	.0066	−.0032	1

TABLE XIII

TOP-3 SELECTED FROM D0, EVALUATED AT D1 VS. A FRESH SELECTION. RANKS ARE OVER ALL 156 QUBITS AT D1.

Policy	selected top-3	$C(q)$ @D1	rank@D1	worst ratio
stale (D0)	q8, q120, q153	.0095/ .0115/ .0139	6/19/32	2.1×
fresh (D1)	q19, q152, q154	.0066/ .0071/ .0071	1/2/3	1.0×

E. Refresh policies

We compare three refresh policies. Let $\text{Jacc}(S, S')$ denote the top-10 Jaccard overlap between two snapshots, and $\Delta C(S, S') = C_{S'}(q^*) - C_S(q^*)$ the absolute change of the stored best qubit’s score.

- 1) **Static (P1)**: never refresh; the initial snapshot S_0 is used for the entire run.
- 2) **Periodic (P2- δ)**: refresh every δ hours regardless of device state ($\delta \in \{6, 12, 24, 48\}$).
- 3) **Adaptive (P3)**: refresh when $\text{Jacc}(S_{\text{stored}}, S_{\text{current}}) < \tau_J$ or $\Delta C(S_{\text{stored}}, S_{\text{current}}) > \tau_C$. The default trigger uses $\tau_J=0.5$, $\tau_C=0.005$.

Re-scoring costs ~ 0.04 ms, so the adaptive policy’s overhead is dominated by the API call for snapshot fetch, not by the staleness check itself. The threshold sweep of §VII-F varies τ_J and τ_C to map the Pareto tradeoff between refresh rate and fidelity.

F. Threshold sweep

To get a *unified, cost-aware* comparison we Monte-Carlo a 30-day runtime (600 jobs) over synthetic drift whose daily step is drawn from the *measured* D0/D1/D2 drift distribution—background $\mathcal{N}(\hat{\mu}, \hat{\sigma})$ robustly estimated from the 156×2 real per-qubit deltas plus a 2% “burst” component (exponential scale 0.15) to reproduce q37-type single-qubit spikes. Every policy runs on the same timelines and reports the same metrics (Table XIV): the adaptive trigger achieves the same fidelity as the daily periodic policy at *under half* the refresh load (0.63 vs. 1.0 refresh/day), and static—which never refreshes—falls below both while missing essentially every drift event.

Sweeping $\tau_J \in [0.2, 0.9]$, $\tau_C \in [0.002, 0.02]$ over the same timelines (Fig. 6): the τ_J term is the robust detector. For $\tau_J \geq 0.4$ fidelity is flat at the plateau for every τ_C ; dropping τ_C to 0.02 or τ_J to 0.2–0.3 relaxes the trigger into missing the burst drift and fidelity breaks below the plateau. The default

TABLE XIV

UNIFIED POLICY COMPARISON, 30 SIMULATED DAYS / 600 JOBS, DRIFT CALIBRATED TO REAL D0/D1/D2 STATISTICS. FIDELITY IS THE FRACTION OF JOBS EXECUTED FROM A STORED SNAPSHOT WHOSE BEST QUBIT STAYS WITHIN TOLERANCE.

Policy	Fidelity	Refresh/day	False	Missed
Static (never)	0.913	0.0	—	117
Periodic 48h	0.908	0.5	0	60
Periodic 24h	0.928	1.0	0	0
Periodic 12h	0.928	2.0	0	0
Periodic 6h	0.928	4.0	0	0
Adaptive (0.5, 0.005)	0.928	0.63	12	41




Fig. 6. Refresh-rate-vs.fidelity frontier. Fidelity saturates once refresh/day $\geq \approx 0.23$; the adaptive operating point (0.63) sits on the plateau, while the periodic schedules spend 1–4 $\times$ more refreshes for no additional fidelity and static falls short of it.

(0.5, 0.005) is at the knee: tightening it further buys no fidelity (fidelity saturates at ~ 0.23 refresh/day) while loosening it sacrifices the plateau. This is the same conclusion as §VII-C stated in operating terms—refresh when accumulated drift approaches the top margin, not on a fixed clock—and it is what makes the choice of $\tau_J=0.5$, $\tau_C=0.005$ rational rather than ad hoc.

G. Where refresh helps: workload and device scale

a) *Workload generalization.*: The refresh *advantage* is not fixed; it grows with how many gates execute on the pinned qubits. Modeling a job’s failure as proportional to the gates on its stored-best qubit (each 2Q gate carrying roughly twice a 1Q gate’s error) and anchoring fresh best $C=0.0066$ and one-day-stale 0.014 (Table XIII): the fresh-minus-stale pass advantage climbs from +0.21 (shallow, sparse) to +0.32 (shallow, dense) and +0.40 (deep, sparse) before both saturate as circuits get so demanding that even a fresh pin fails (Table XV, lower block). The read is exactly the opposite of counter-intuitive: the adaptive-refresh gain is largest on the deep, 2Q-heavy workloads that utility-scale execution actually cares about—refresh is not a small-circuit optimization.

b) *Device and layout scale.*: The refresh benefit is also k - and topology-dependent (Table XV, upper block). Searching the real 156-qubit heavy-hex graph for connected k -layouts under the D0 and D1 snapshots shows the stale quality penalty (a D0-optimal layout re-evaluated at D1 minus a fresh D1 pick) is concentrated at *intermediate* k : zero to a few points at $k \leq 8$, peaking at $k=16$ (+0.045) and $k=32$ (+0.026), then shrinking to zero by $k \geq 100$. At large k the layout is forced to include the full device, so its mean readout approaches the global average regardless of freshness—qubit quality stops being the lever.

TABLE XV

WHERE THE REFRESH BENEFIT LIVES. UPPER: STALE QUALITY PENALTY = (D0-OPTIMAL LAYOUT COST AT D1) – (FRESH D1 PICK COST) OVER THE REAL HEAVY-HEX MAP, PLUS THE CONNECTIVITY SHARE OF LAYOUT-COST VARIANCE. LOWER: WORKLOAD PASS-RATE OF A FRESH VS. A ONE-DAY-STALE PIN AS CIRCUIT DEPTH AND 2Q DENSITY GROW (fresh $C=0.0066$, stale $C=0.014$).

k	stale penalty	conn. share
4	+0.002	0.00
16	+0.045	0.00
32	+0.026	0.02
64	+0.006	0.11
100	0.000	0.35
128	0.000	0.77
workload	pass fresh	pass stale
shallow, sparse	0.82	0.61
shallow, dense	0.71	0.39
deep, sparse	0.45	0.05
deep, dense	0.13	0.05

The same k -sweep shows the share of layout-cost variance explained by connectivity (induced 2Q/boundary cost) rising monotonically: ~ 0 at $k \leq 16$, 0.11 at $k=64$, 0.35 at $k=100$, 0.77 at $k=128$ —the crossover where connectivity dominates qubit quality sits near the device’s physical scale. Thus refresh is most valuable in the operating regime $8 \lesssim k \lesssim 64$ where a runtime both has placement freedom *and* drift alters which free pick is best; for near-device-scale k the bottleneck is connectivity, which refresh cannot change.

VIII. COMPOSING THE SOFTWARE AND HARDWARE PATHS

A. Reliability calculus of the orchestration layer

Because the orchestrator’s outcome paths are mutually exclusive, end-to-end reliability decomposes exactly as the chain rule over its sequential gates

$$R_{\text{end}} = P(\text{first_pass}) + P(\text{enter heal})\rho_{\text{heal}} + P(\text{enter fallback})\rho_{\text{fb}}, \quad (5)$$

with no independence assumption needed for the identity. The healing stage is modeled from the implementation’s control flow (at most $K=3$ generation calls stopping at first success; healing budget $H=4$ with abort-on-None semantics): with per-call probabilities g (correct), r (broken), f (unavailable; $g+r+f=1$),

$$G = 1 + f + f^2, \quad \rho_{\text{heal}} = g(1 + r + r^2 + r^3), \quad (6)$$

$$R_{\text{end}} = gG + rG\rho_{\text{heal}} + f^3c, \quad (7)$$

with c the fallback coverage. This model was validated cell-by-cell against all 23 testable outcome cells of the ablation and availability-stress grids: *every* prediction falls inside the observed Wilson 95% interval (Table XVI shows the decomposition and stress cells). At $q=0.5$ the structural prediction $\rho_h=81.2\%$ sits 0.3 pp from the measured 81.5%, while the naive independence baseline $1 - (1 - q)^4 = 93.8\%$ overshoots by 12.4 pp; the model also separates the paths (first-pass 55.5% predicted vs. 54.4% measured; healed path 36.1% vs. 37.1%). The correction matters: the naive-measured gap reported in early drafts is a retry-budget truncation effect, not stochastic

TABLE XVI

RELIABILITY-COMPOSITION VALIDATION (450 TRIALS PER CELL). ALL PREDICTED VALUES INSIDE THE WILSON 95% INTERVAL; + = WITHIN CI.

Cell	Pred	Obs	CI	✓
<i>Decomposition ρ_h at $q=0.5$</i>				
structural	0.812	0.815	[0.756, 0.862]	+
naive indep	0.938	—	gap 12.4 pp	—
<i>Decomposition at $q=0.7$</i>				
structural	0.874	0.914	[0.839, 0.956]	+
naive indep	0.992	—	gap 7.8 pp	—
<i>Availability stress E06</i>				
S2 fr=0.3	0.973	0.960	[0.938, 0.975]	+
S2 fr=0.5	0.875	0.864	[0.830, 0.893]	+
S3 fr=0.3	0.979	0.969	[0.948, 0.981]	+
S3 fr=0.5	0.903	0.896	[0.864, 0.921]	+

TABLE XVII

DESIGN-SPACE RANKING OF ORCHESTRATION FIXES AT ($q=0.5$, $\text{fr}=0.1$).

Variant	ρ_h	Δ (pp)
baseline (abort-on-None, $H=4$)	0.812	—
None wastes a slot (no abort), $H=4$	0.938	+12.5
budget $H=8$ (abort retained)	0.833	+2.1
healer quality $q_h=0.9$	0.900	+8.8

correlation—under abort-on-None, unavailability both wastes heal slots and terminates the loop. The model also ranks fixes before code is written (Table XVII): treating None as wasted-but-nonterminal recovers the entire impression (+12.5 pp) at near-zero engineering cost; doubling the budget to $H=8$ buys only +2.1 pp; raising the healer quality to $q_h=0.9$ buys +8.8 pp. The fallback’s effective coverage is $c=2/9$ (nominal keyword coverage over a 9-intent benchmark: seven of nine contain apostrophes that break the template interpolation)—a stated limitation. We include this model because it is the *software* side of the composition; the paper’s hardware claims are placement and drift, not LLM quality, which we keep as a controlled parameter.

B. End-to-end placement batch

To test the composable value of placement *without* the staleness confound, we ran the same 6-qubit angle suite (30 parameterizations, 4096 shots, tolerance 0.05 on all six per-bit deviations) through **full** (greedy chain pinned from the freshly fetched snapshot) and **ablated** (free placement), back-to-back on 2026-08-31; Table XIX also shows the earlier identical batch pinned from D0 to separate temporal from placement effects.

Two effects separate cleanly. First, *temporality*: refreshing pins lifts *both* arms by +33/37 pp, including the ablated arm that never saw a pin—device state at run time dominates. Second, *placement shows no measured advantage* in either batch (−3.3 and −6.7 pp, both non-significant): at $k=6$ with fresh pins, free transpiler placement already matches our best chain. The failing-bit anatomy is instructive. In the refresh-full arm the 12 out-of-tolerance bit-events concentrate on late-chain qubits q11 (4), q13 (4), q12 (1)—the chain’s forced inclusion problem again, now visible even with fresh pins. In

TABLE XVIII

OUT-OF-TOLERANCE BIT-EVENTS PER END-TO-END ARM AND THEIR CONCENTRATION. FROZEN BATCHES USE D0 PINS; REFRESH BATCHES PIN FROM THE 08-31 SNAPSHOT.

Arm	events	dominant qubits (events)
frozen full	20	q4 (16), q7 (3)
frozen ablated	18	q4 (16), q1 (1), q3 (1)
refresh full	12	q11 (4), q13 (4), q12 (1)
refresh ablated	7	q4 (5), q3 (2)

TABLE XIX

6-QUBIT ANGLE BATCHES ($n=30/\text{ARM}$, 4096 SHOTS). PASS = ALL SIX PER-BIT DEVIATIONS ≤ 0.05 . WILSON 95% CI; TWO-PROPORTION p FOR FULL VS. ABLATED WITHIN A BATCH.

Batch	Full (pinned)	Ablated (free)	p
Frozen pins (D0)	36.7% [21.9,54.5]	40.0% [24.6,57.7]	0.79
Refresh pins (runtime)	70.0% [52.1,83.3]	76.7% [59.1,88.2]	0.56
Δ	+33.3 pp	+36.7 pp	—

the refresh-ablated arm the 7 events sits on q4 (5) and q3 (2): q4 was rank 4 at D0 but its C rises $0.0078 \rightarrow 0.0173$ by D1/D2 (rank 52), yet the free layout, which cannot see the trend, still lands on it repeatedly—drift *inside* the layout choice, unmitigated by any pin. Table XVIII collects the four arms. This is the honest scope of §V: placement is decisive between engineered extremes (q98 vs. q37: 0.028 vs. 0.213) and sets the reliability *floor* at $k \geq 9$ under stale pins, but its average benefit against free placement at small k is within temporal noise. The full pipeline (including the LLM orchestration leg) measured 54/54 completed by the full arm vs. 47/54 (87.0%) by the ablated one ($z=2.74$, one-sided $p=0.003$), consistent with the model’s 97.3% prediction; we report the earlier $n=54$ frozen composition (74.1% vs. 66.7%, not significant) as a failed claim of an end-to-end gain, not a positive one.

IX. DISCUSSION

A. What the measurements do and do not say

Three claims survive the data as stated: (i) encoding fidelity is placement-dominated—the engineered extremes differ by $\sim 8\times$ (and a calibration-predictable affine channel describes the distortion); (ii) device calibration ordering moves on a daily cycle to a degree ($\rho=0.58$, top-10 Jaccard 0.25, stale worst pin $2.1\times$ worse) that makes selection age the first-order risk, and refresh is cheap; (iii) simulator-proxy training granted $J(G)$ an out-of-sample generalization ($\rho=0.82$) that real labels refute ($\rho=0.13$). Claims that do *not* survive, herewith withdrawn or bounded: the first run’s perfect monotone ranking; the top-half rule; and any use of proxy-fitted $J(G)$ at $k \geq 11$.

B. Operational consequences for quantum-cloud runtimes

The marginal cost of a calibration fetch and a scoring pass is effectively zero, so the expensive resource is *stale trust*, not compute. A runtime should (a) timestamp layouts against their snapshot, (b) refresh on the trigger rather than a fixed wall-clock, (c) train scheduling models exclusively on real,

TABLE XX

MEASURED FAILURE MODES AND THE RESPONSE EACH SUPPORTS. PRIORITY IS BY EVIDENCE STRENGTH, NOT EASE.

Measured failure	Evidence	Response
unpinned layout misses tolerance	§V: 0.392 vs. 0.028	score-pinned <code>initial_layout</code>
stale snapshot invalidates selection	§VII: stale best pin $2.1\times$ worse	trigger-based refresh; snapshot age
proxy labels fool $J(G)$	§VI: OOS ρ 0.82 $\rightarrow$ 0.13	refit on real labels; ban $k \geq 11$ extrapolation
chain forced-inclusion at large k	§VI/§VIII: q11,q13, q67 fail	connectivity-vs- C trade-off (open)
abort-on-None truncates heal loop	§VIII-A: +12.5 pp fix	wasted-slot semantics

refresh-pinned labels, and (d) expose snapshot age as a first-class job attribute—all implementable without changing the quantum programming model, which is exactly what UniMind does. Where math precedes implementation, the reliability calculus ranks the fixes (§VIII-A), and the endowment to port is a calibration snapshot plus a scoring pass, neither of which is exotic.

C. Design synthesis

Table XX maps each measured failure mode to the engineering response the data supports; the columns are ordered by measured cost so a runtime can prioritize. The two lines a naive reading hides—“mitigation helps” and “ $J(G)$ ranks well”—are the ones our data bounds, and both bounds are the point of the paper.

D. Threats to validity

Measurement noise: 4096–8192 shots put per-bit deviations on a floor of ≈ 0.02 ; the 0.05 tolerance is coarse but was fixed before data release. **Drift generality:** D0–D2 are three snapshots, one device, one topology family (heavy-hex); the value is the protocol, not the constants. **Placement null power:** $n=30/\text{arm}$ can detect only large effects; we report point estimates and CIs rather than asserting equivalence. **Selection fairness:** the greedy chain maximizes $C(q)$ -locality and at large k is forced to include high-error boundary qubits; a connectivity-vs- C trade-off could change §VI, and we report that failure as motivation. **Single-model LLM layer:** the orchestration calculus holds under i.i.d. mock semantics; real-model autocorrelation is explicitly outside its scope. **Utility-vs-layout coupling:** the $J(G)$ features are computed on the transpiled layout; errors in E_{2q} propagate to the weight fit but not to the device measurements, which are layout-real.

X. CONCLUSION

We presented a measurement-driven account of calibration-aware scheduling for quantum-accelerated AI middleware on a 156-qubit device: a score-only placement pipeline with

measured scope, daily drift dynamics with a working refresh trigger, an exact reliability composition of the software path, an end-to-end test showing that at our scale refresh dominates selection while selection still sets the placement floor, and the explicit negative results—the non-replicating monotone ranking, the absence of a measured placement advantage at $k=6$, and the collapse of the proxy-trained utility under real labels. The command line for future work is set by the failures: nightly D2–D5 drift accumulation, refresh-conditional placement policy ($C(q)$ -vs-connectivity trade-off, including the q11/q13 late-chain failures observed under fresh pins), and active-learning refits of $J(G)$ on real labels across devices. The open repository makes every number traceable to a job or snapshot.

REFERENCES

- [1] J. Preskill, “Quantum computing in the NISQ era and beyond,” *Quantum*, vol. 2, p. 79, 2018.
- [2] F. Arute *et al.*, “Quantum supremacy using a programmable superconducting processor,” *Nature*, vol. 574, pp. 505–510, 2019.
- [3] Y. Kim *et al.*, “Evidence for the utility of quantum computing before fault tolerance,” *Nature*, vol. 618, pp. 500–505, 2023.
- [4] V. Bergholm *et al.*, “PennyLane: Automatic differentiation of hybrid quantum-classical computations,” *arXiv:1811.04968*, 2018.
- [5] P. Murali, J. M. Baker, A. Javadi-Abhari, F. T. Chong, and M. Martonosi, “Noise-adaptive compiler mappings for noisy intermediate-scale quantum computers,” in *Proc. ASPLOS*, 2019, pp. 1015–1029.
- [6] G. Li, Y. Ding, and Y. Xie, “Tackling the qubit mapping problem for NISQ-era quantum devices,” in *Proc. ASPLOS*, 2019, pp. 1001–1014.
- [7] B. Tan and J. Cong, “Optimal layout synthesis for quantum computing,” in *Proc. ICCAD*, 2020, pp. 1–9.
- [8] L. Lao and D. E. Browne, “Timing and resource-aware mapping of quantum circuits to heterogeneous MWF-gate superconducting quantum processors,” *ACM Trans. Quantum Comput.*, vol. 3, no. 3, 2022.
- [9] K. Temme, S. Bravyi, and J. M. Gambetta, “Error mitigation for short-depth quantum circuits,” *Phys. Rev. Lett.*, vol. 119, p. 180509, 2017.
- [10] E. van den Berg, Z. K. Mineev, A. Kandala, and K. Temme, “Probabilistic error cancellation with sparse Pauli–Lindblad models on noisy quantum processors,” *Nature Physics*, vol. 19, pp. 1116–1121, 2023.
- [11] P. D. Nation *et al.*, “Scalable mitigation of measurement errors on quantum computers,” *PRX Quantum*, vol. 5, p. 010326, 2024.
- [12] S. Bravyi, S. Shaydulin, S. Hu, and K. Temme, “Mitigating measurement errors in multiqubit experiments,” *Phys. Rev. Lett.*, vol. 127, p. 200503, 2021.
- [13] N. Kanazawa *et al.*, “Quantum error mitigation,” *arXiv:2210.00921*, 2022.
- [14] R. LaRose *et al.*, “Mitiq: a software package for error mitigation on noisy quantum computers,” *Quantum*, vol. 6, p. 749, 2022.
- [15] L. Viola, E. Knill, and S. Lloyd, “Dynamical decoupling of open quantum systems,” *Phys. Rev. Lett.*, vol. 82, pp. 2417–2421, 1999.
- [16] “Qiskit Runtime primitives,” IBM Quantum documentation, [Online]. Available: <https://docs.quantum.ibm.com/api/qiskit-ibm-runtime>
- [17] “IBM Quantum Platform calibration data,” [Online]. Available: <https://quantum.ibm.com/services/resources>
- [18] A. Javadi-Abhari *et al.*, “Quantum computing with Qiskit,” *arXiv:2405.08810*, 2024.
- [19] NVIDIA, “cuQuantum SDK,” [Online]. Available: <https://developer.nvidia.com/cuquantum-sdk>
- [20] Amazon, “Amazon Braket,” [Online]. Available: <https://aws.amazon.com/braket>
- [21] Microsoft, “Azure Quantum,” [Online]. Available: <https://azure.microsoft.com/quantum>
- [22] M. Schuld and F. Petruccione, *Supervised Learning with Quantum Computers*. Cham: Springer, 2018.
- [23] V. Havlíček *et al.*, “Supervised learning with quantum-enhanced feature spaces,” *Nature*, vol. 567, pp. 209–212, 2019.
- [24] Y. X. Tan, “UniMind: a middleware framework for bridging classical AI workloads and quantum computing backends,” tech. report v2.5, 2026 (companion repo `yxtan5022-maker/unimind-v2`).
- [25] UniMind placement/drift datasets (2026-08), `data/` in `yxtan5022-maker/unimind-v2`.

APPENDIX

A. Per-row results of the $k=9$ –18 suite

The 14 circuits (seven sizes $\times$ two seeded patterns) with greedy-chain placement, 4096 shots, tolerance 0.05. Worst qubit is the failing or most-off bit; ranks are over all 156 qubits at D1.

TABLE XXI
PER-ROW MAX DEVIATION UNDER FRESH (D1/D2) AND STALE (D0) PINS.

k	pat.	fresh	worst@D1	stale	worst@D1
9	0	.064	q11 (71)	.151	q4 (52)
9	1	.056	q11 (71)	.045	q24 (118)
10	0	.053	q14 (56)	.118	q23 (136)
10	1	.035	q13 (29)	.088	q23 (136)
11	0	.081	q11 (71)	.049	q4 (52)
11	1	.042	q11 (71)	.145	q4 (52)
12	0	.041	q10 (7)	.067	q4 (52)
12	1	.038	q9 (74)	.107	q4 (52)
14	0	.037	q13 (29)	.183	q46 (126)
14	1	.077	q11 (71)	.096	q23 (136)
16	0	.040	q4 (52)	.265	q67 (132)
16	1	.085	q4 (52)	.239	q67 (132)
18	0	.089	q11 (71)	.332	q67 (132)
18	1	.047	q13 (29)	.150	q67 (132)

B. Data and reproducibility

Every table is backed by traced data in the repository: snapshots, E2E batches, real $J(G)$ labels, and the analysis scripts listed below.

```
data/drift/calib_2026-{0829,0830,08-31}.json
data/e2e/e2e_angle_{full,ablated}*.json
data/jgpu/j_real*.json
analysis/*.py + analysis/results/*.json
analysis/results/refresh_policy_analysis.json
analysis/results/power_analysis.json
analysis/results/jg_failure_analysis.json
```

QPU job IDs appear in the file headers: daadeocjbipc73ffq83g (refresh-full E2E), daadeosjbipc73ffq840 (refresh-ablated E2E), and daadgmmrbfbs73cijmbg (refresh-pinned $J(G)$ labels). Transpiler seed 42 everywhere; shots per experiment as stated.

C. Calibration-weighted routing benchmark

Interactive (local-simulator, zero-quota) comparison on snapshot 2026-08-29 (156 qubits, 176 couplers; RY–CX–RY, opt=1, seed 42). UniMind computes a $C(q) + s_x$ -weighted SABRE candidate set and keeps the better of weighted and global SABRE, guaranteeing no regression for $k \geq 8$ (violation check: none).

TABLE XXII

PLACEMENT COST OF THE FIVE STRATEGIES (SUMMARY ROWS; FULL PER- k TABLES IN THE REPO).

k	2	4	6	8	10	16
Random SWAPs	8	42	68	105	118	207
UniMind SWAPs	0	3	8	20	20	50
Δ SWAP vs. random	-100%	-93%	-88%	-81%	-83%	-76%
Random 2q err	.118	.484	.666	.821	.847	.979
UniMind 2q err	.002	.049	.065	.133	.181	.353
Δ vs. Default	= 0	= 0	= 0	= 0	= 0	= 0